

# Processes and Threads

1 UNSW

1

## Learning Outcomes

- An understanding of fundamental concepts of processes and threads

2 UNSW

2

## Major Requirements of an Operating System

- Interleave the execution of several processes to maximize processor utilization while providing reasonable response time
- Allocate resources to processes
- Support interprocess communication and user creation of processes

3 UNSW

3

## Processes and Threads

- Processes:
  - Also called a task or job
  - Execution of an individual program
  - "Owner" of resources allocated for program execution
  - Encompasses one or more threads
- Threads:
  - Unit of execution
  - Can be traced
    - list the sequence of instructions that execute
  - Belongs to a process
    - Executes within it.

4 UNSW

4

Execution snapshot of three single-threaded processes (No Virtual Memory)

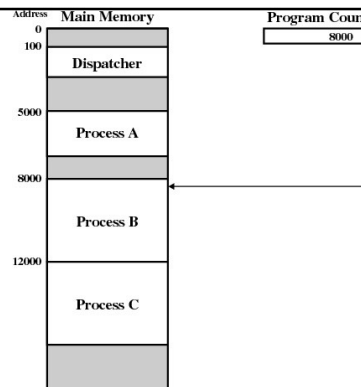


Figure 3.1 Snapshot of Example Execution (Figure 3.1 at Instruction Cycle 13)

5

## Logical Execution Trace

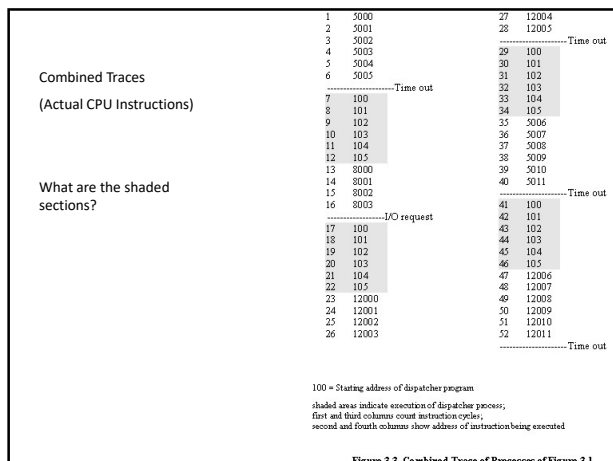
5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

(a) Trace of Process A (b) Trace of Process B (c) Trace of Process C

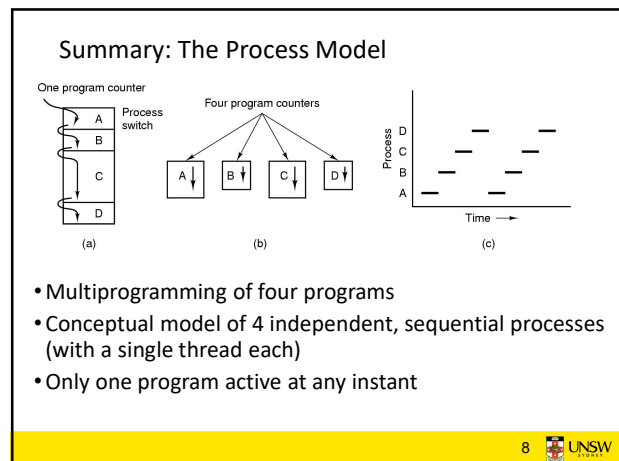
5000 = Starting address of program of Process A  
8000 = Starting address of program of Process B  
12000 = Starting address of program of Process C

Figure 3.2 Traces of Processes of Figure 3.1


6



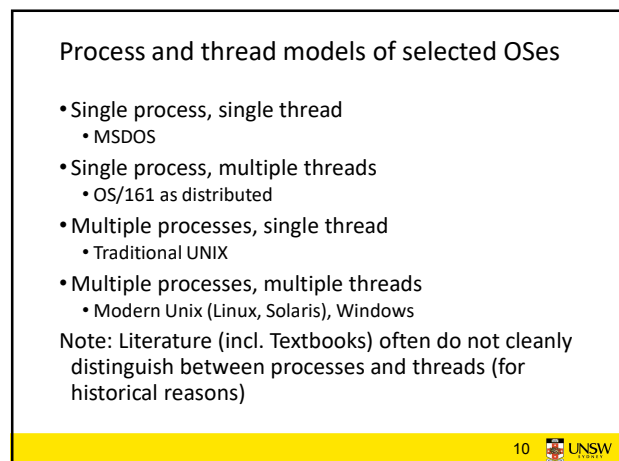
7



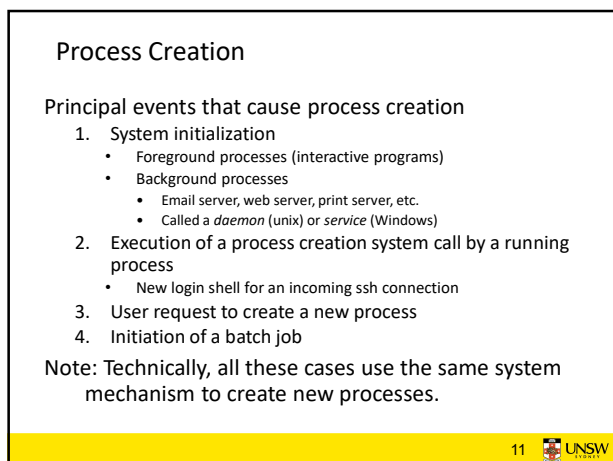
8



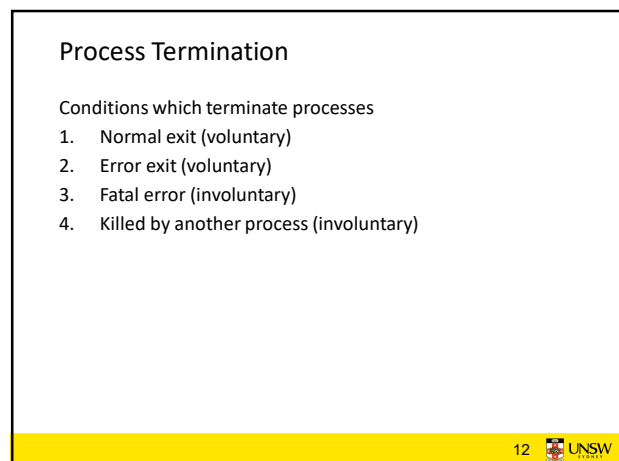
9



10




11



12

## Implementation of Processes

- A processes' information is stored in a *process control block* (PCB)
- The PCBs form a *process table*
  - Reality can be more complex (hashing, chaining, allocation bitmaps,...)



13 UNSW

13

## Implementation of Processes


Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment Pointer to data segment Pointer to stack segment	Root directory Working directory File descriptors User ID Group ID

Example fields of a process table entry

14 UNSW

14

## Process/Thread States



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

- Possible process/thread states
  - running
  - blocked
  - ready
- Transitions between states shown

15 UNSW

15

## Some Transition Causing Events

### Running → Ready

- Voluntary `yield()`
- End of timeslice

### Running → Blocked

- Waiting for input
  - File, network,
- Waiting for a timer (alarm signal)
- Waiting for a resource to become available

16 UNSW

16


## Scheduler

- Sometimes also called the *dispatcher*
  - The literature is also a little inconsistent on with terminology.
- Has to choose a *Ready* process to run
  - How?
  - It is inefficient to search through all processes

17 UNSW

17

## The Ready Queue



18 UNSW


18

## What about blocked processes?


- When an *unblocking* event occurs, we also wish to avoid scanning all processes to select one to make *Ready*

19 UNSW

## Using Two Queues




20 UNSW



21

## Threads The Thread Model



22 UNSW

## The Thread Model – Separating execution from the environment.


### Per process items

Address space  
Global variables  
Open files  
Child processes  
Pending alarms  
Signals and signal handlers  
Accounting information

### Per thread items

Program counter  
Registers  
Stack  
State

- Items shared by all threads in a process
- Items private to each thread



23

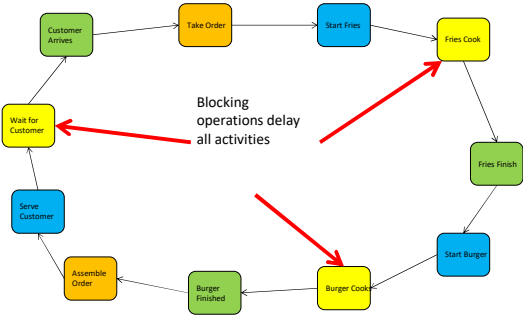
## Threads Analogy



The Hamburger Restaurant

24 UNSW

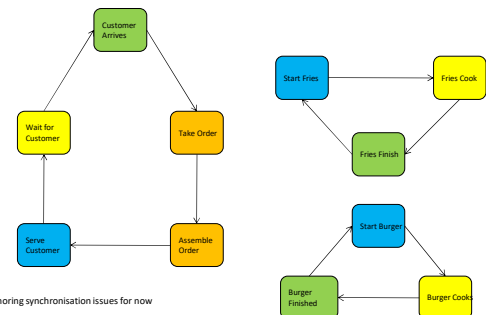
### Single-Threaded Restaurant



25 UNSW

25

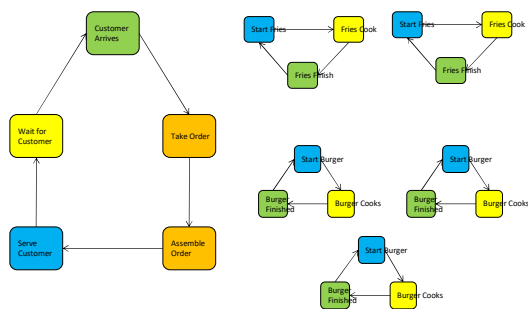
### Multithreaded Restaurant



26 UNSW

26

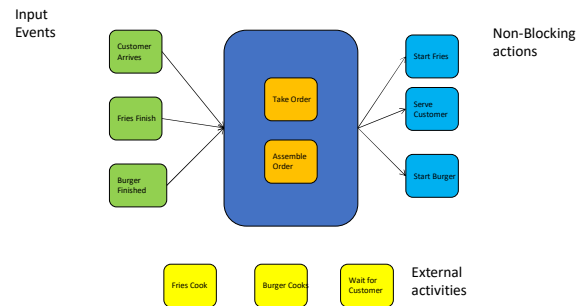
### Multithreaded Restaurant with more worker threads



27 UNSW

27

### Finite-State Machine Model (Event-based model)




28 UNSW

28

### Observation: Computation State

#### Thread Model



- State implicitly stored on the stack.

#### Finite State (Event) Model




- State explicitly managed by program

29 UNSW

29

### The Thread Model



Each thread has its own stack

30 UNSW

30


## Thread Model

- Local variables are per thread
  - Allocated on the stack
- Global variables are shared between all threads
  - Allocated in data section
  - Concurrency control is an issue
- Dynamically allocated memory (malloc) can be global or local
  - Program defined (the pointer can be global or local)

31 UNSW

31

## Thread Usage




A word processor with three threads

32 UNSW

32

## Thread Usage



A multithreaded Web server

33 UNSW

33

## Thread Usage

```

while (TRUE) {
    get_next_request(&buf);
    handoff_work(&buf);
}
(a)

while (TRUE) {
    wait_for_work(&buf);
    look_for_page_in_cache(&buf, &page);
    if (page_not_in_cache(&page))
        read_page_from_disk(&buf, &page);
    return_page(&page);
}
(b)
    
```

- Rough outline of code for previous slide
  - (a) Dispatcher thread
  - (b) Worker thread – can overlap disk I/O with execution of other threads

34 UNSW

34

## Thread Usage

Model	Characteristics
Threads	Parallelism, blocking system calls
Single-threaded process	No parallelism, blocking system calls
Finite-state machine	Parallelism, nonblocking system calls, interrupts

Three ways to construct a server

35 UNSW

35

## Summarising “Why Threads?”

- Simpler to program than a state machine
- Less resources are associated with them than a complete process
  - Cheaper to create and destroy
  - Shares resources (especially memory) between them
- Performance: Threads waiting for I/O can be overlapped with computing threads
  - Note if all threads are *compute bound*, then there is no performance improvement (on a uniprocessor)
- Threads can take advantage of the parallelism available on machines with more than one CPU (multiprocessor)

36 UNSW

36