



COMP3231/9201/3891/9283

Operating Systems 2021/T1

UNSW

Administration

- [Notices](#)
- [Course Outline](#)
- [UNSW Timetable](#)
- [Consultations](#)
- [Survey Results!!](#)

Work

- [Lectures](#)
- [Tutorials](#)

Support

- [Ed Forums](#)
- [Wiki](#)

Assignments

- [Submission Guide](#)
- [Assignment 0 Warm-up](#)

Resources

OS/161

- [General](#)
- [Man Pages](#)
- [Sys161 Pages](#)

C coding

- [Info Sheet](#)

Debugging

- [Learn Debugging](#)

General

- ["Hardware" Guide](#)
- [R3000 Reference Manual](#)
- [Intro. to Prog. Threads](#)

Previous years

- [2020 T2](#)
- [2020 T1](#)
- [2019 T1](#)
- [2018 S1](#)
- [2017 S1](#)
- [2016 S1](#)
- [2015 S1](#)
- [2014 S1](#)
- [2013 S1](#)
- [2012 S1](#)
- [2011 S1](#)
- [2010 S1](#)
- [2009 S1](#)
- [2008 S1](#)
- [2007 S1](#)
- [2006 S1](#)
- [2005 S2](#)
- [2005 S1](#)
- [2004 S2](#)
- [2004 S1](#)

Tutorial Week 7

Questions

Files and file systems

1. Why does Linux pre-allocate up to 8 blocks on a write to a file.

2. Linux uses a *buffer cache* to improve performance. What is the drawback of such a cache? In what scenario is it problematic? What alternative would be more appropriate where a buffer cache is inappropriate?

3. What is the structure of the contents of a directory? Does it contain attributes such as creation times of files? If not, where might this information be stored?

4. The Unix inode structure contains a reference count. What is the reference count for? Why can't we just remove the inode without checking the reference count when a file is deleted?

5. Inode-based filesystems typically divide a file system partition into *block groups*. Each block group consists of a number of contiguous physical disk blocks. Inodes for a given block group are stored in the same physical location as the block groups. What are the advantages of this scheme? Are there any disadvantages?

6. Assume an inode with 10 direct blocks, as well as single, double and triple indirect block pointers. Taking into account creation and accounting of the indirect blocks themselves,

Staff

- [Kevin Elphinstone \(LiC\)](#)
- TBD (Admin)

Grievances

- [Student Reps](#)



what is the largest possible number of block reads and writes in order to:

- Read 1 byte
- Write 1 byte

Assume the inode is cached in memory.

- Assume you have an inode-based filesystem. The filesystem has 512 byte blocks. Each inode has 10 direct, 1 single indirect, 1 double indirect, and 1 triple indirect block pointer. Block pointers are 4 bytes each. Assume the inode and any block free list is always in memory. Blocks are not cached.
 - What is the maximum file size that can be stored before
 - the single indirect pointer is needed?
 - the double indirect pointer is needed?
 - the triple indirect pointer is needed?
 - What is the maximum file size supported?
 - What is the number of disk block reads required to read 1 byte from a file
 - in the best case?
 - in the worst case?
 - What is the number of disk block reads and writes required to write 1 byte to a file
 - in the best case?
 - in the worst case?

- A typical UNIX inode stores both the file's size and the number of blocks currently used to store the file. Why store both? Shouldn't blocks = size / block size?

- How can deleting a file leave a inode-based file system (like ext2fs in Linux) inconsistent in the presence of a power failure.

- How does adding journalling to a file system avoid corruption in the presence of unexpected power failures.

Page last modified: 4:49pm on Saturday, 23rd of May, 2020

[Print Version](#)

CRICOS Provider Number: 00098G